



# mongoDB

## CS572 Modern Web Applications Programming Maharishi University of Management Department of Computer Science Assistant Professor Asaad Saad

Except where otherwise noted, the contents of this document are Copyrighted. All rights reserved. Any redistribution, reproduction, transmission, or storage of part or all of the contents in any form is prohibited without the author's expressed written permission. Slides have been modified for Maharishi University of Management Computer Science course CS572. Credit goes to: Andrew Erlichson, Shaun Verch, Richard Kreuter.

1

- OUTLINE
- 🔍
-  1. ---
  -  2. Maharishi University of Management - Fairfield, Iowa
  -  3. CRUD in MongoDB
  -  4. Searching an Array
  -  5. Search with \$elemMatch
  -  6. \$elemMatch
  -  7. Searching an Object
  -  8. Update Methods
  -  9. Field Update Operators





# CRUD in MongoDB

There is no special SQL language to perform CRUD in MongoDB.

Many CRUD operations exist as methods/functions on objects in programming language API, NOT as separate language.

CRUD	MongoDB	SQL
Read	find()	select
Create	insert()	insert
Update	update()	update
Delete	remove()	delete

3

- OUTLINE
- 🔍
- 1. ---
  - 2. Maharishi University of Management - Fairfield, Iowa
  - 3. CRUD in MongoDB**
  - 4. Searching an Array
  - 5. Search with \$elemMatch
  - 6. \$elemMatch
  - 7. Searching an Object
  - 8. Update Methods
  - 9. Field Update Operators



# Searching an Array

```
// { _id: 1, courses: [ "CS472", "CS572", "CS435" ] }
// find all documents where courses value contains "CS572"
db.col.find({ courses: "CS572" })

// find all documents where courses value contains "CS472" or "CS572"
db.col.find({ courses: { $in: ["CS572", "CS472"] } })

// find all documents where courses value contains "CS472" and "CS572"
db.col.find({ courses: { $all: [ "CS572" , "CS472" ] } })
```

The **\$in** operator selects the documents where the value of a field is an array that contains at any order any (at least one) of the specified elements

The **\$all** operator selects the documents where the value of a field is an array that contains at any order all the specified elements

4

## OUTLINE



1. ---



2. Maharishi University of Management - Fairfield, Iowa



3. CRUD in MongoDB



4. Searching an Array



5. Search with \$elemMatch



6. \$elemMatch



7. Searching an Object



8. Update Methods



9. Field Update Operators





# Search with \$elemMatch

```
{ _id: 1, results: [ 1, 2, 5, 10 ] }
{ _id: 2, results: [ 5, 8, 9, 10 ] }
{ _id: 3, results: [ 10, 11, 12 ] }
```

```
db.test.find( { results: { $all: [ 5, 10 ] } } )
// scan results for value x times (accepts only value, no operators)
// All values must exist to return the document
// scan if it has 5 AND scan again if it has 10 = returns _id: 1
```

```
db.test.find( { results: { $elemMatch: { $gt: 5, $lt: 10 } } } )
// scan results array once: check if there is one value matches the condition
// _id: 2
```

5

- OUTLINE
- Search...
1. ---
  2. Maharishi University of Management - Fairfield, Iowa
  3. CRUD in MongoDB
  4. Searching an Array
  5. Search with \$elemMatch
  6. \$elemMatch
  7. Searching an Object
  8. Update Methods
  9. Field Update Operators



# \$elemMatch

The **\$elemMatch** operator matches documents that contain an array field with at least one **element** that matches **ALL** the specified query criteria.

```
{ _id: 1, results: [ 82, 85, 88 ] }
{ _id: 2, results: [ 75, 88, 89 ] }
```

```
db.scores.find( { results: { $elemMatch: { $gte: 80, $lt: 85 } } } ) // _id: 1
{ _id: 1, results: [ { product: "abc", score: 10 }, { product: "xyz", score: 5 } ] }
{ _id: 2, results: [ { product: "abc", score: 8 }, { product: "xyz", score: 7 } ] }
{ _id: 3, results: [ { product: "abc", score: 7 }, { product: "xyz", score: 8 } ] }

db.survey.find( { results: { $elemMatch: { product: "xyz", score: { $gte: 8 } } } } ) // _id: 3
```

6

## OUTLINE



1. ---



2. Maharishi University of Management - Fairfield, Iowa



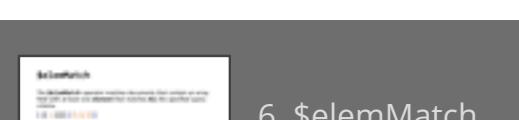
3. CRUD in MongoDB



4. Searching an Array



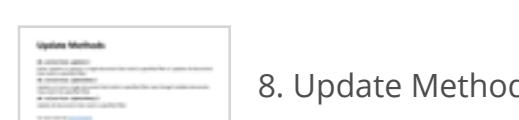
5. Search with \$elemMatch



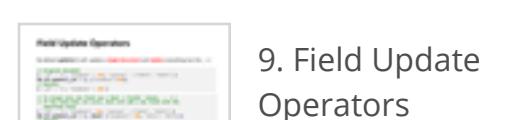
6. \$elemMatch



7. Searching an Object



8. Update Methods



9. Field Update Operators



# Searching an Object

```
{_id: 1, email: { work: "work@mum.edu", personal: "personal@gmail.com" } }
```

// nothing will be returned

```
db.col.find({ email: { work: "work@mum.edu" } })
db.col.find({ email: { personal: "personal@gmail.com" } })
db.col.find({ email: { personal: "personal@gmail.com", work: "work@mum.edu" } })
```

// will work

```
db.col.find({ email: { work: "work@mum.edu", personal: "personal@gmail.com" } })
```

// how to search for one key only?

```
db.col.find({ "email.work": "work@mum.edu" })
db.col.find({ "email.personal": "personal@gmail.com" })
```

7

## OUTLINE



1. ---



2. Maharishi University of Management - Fairfield, Iowa



3. CRUD in MongoDB



4. Searching an Array



5. Search with \$elemMatch



6. \$elemMatch



7. Searching an Object



8. Update Methods



9. Field Update Operators





# Update Methods

## **db.collection.update()**

Either updates or replaces a single document that match a specified filter or updates all documents that match a specified filter.

## **db.collection.updateOne()**

Updates at most a single document that match a specified filter even though multiple documents may match the specified filter.

## **db.collection.updateMany()**

Update all documents that match a specified filter.

For more check the [documentation](#)

8

## OUTLINE



1. ---



2. Maharishi University of Management - Fairfield, Iowa



3. CRUD in MongoDB



4. Searching an Array



5. Search with \$elemMatch



6. \$elemMatch



7. Searching an Object



8. Update Methods



9. Field Update Operators





# Field Update Operators

By default `update()` will: update a **single document** and **replace** everything but the `_id`

```
// Original document
{ "_id" : "1", "students" : 250, "courses" : ["CS572", "CS472"] }
db.col.update({_id:"1"}, {"students":500})
// Results
{ "_id" : "1", "students" : 500 }
```

```
// To target only one field use { $set: { field1: value1, ... } }
// If the field does not exist, $set will add a new field with the
// specified value
{ "_id" : "1", "students": 250, "courses" : ["CS572", "CS472"] }
db.col.update({_id:"1"}, {$set: {"students": 500, "entry": "Oct" } })
// Results
{ "_id" : "1", "students": 500, "courses" : ["CS572", "CS472"] , "entry":"Oct" }
```

9

## OUTLINE



1. ---

2. Maharishi University of Management - Fairfield, Iowa

3. CRUD in MongoDB

4. Searching an Array

5. Search with \$elemMatch

6. \$elemMatch

7. Searching an Object

8. Update Methods

9. Field Update Operators





# Field Update Operators

```
// Original document
{ "_id" : "1", "students" : 250, "courses" : ["CS572", "CS472"] }
db.col.update({_id:"2"}, {"students":500}, {upsert: true})
// Results
{ "_id" : "1", "students" : 250, "courses" : ["CS572", "CS472"] }
{ "_id" : "2", "students" : 500 }
```

```
// update all docs to have one more field (city: Fairfield)
{ "_id" : "1", "Dept" : "CompPro" }
{ "_id" : "2", "Dept" : "SustainableLiving" }
db.col.update({}, {$set: {"city":"Fairfield"}}, {multi: true})
// Results
{ "_id" : "1", "Dept" : "CompPro", "city":"Fairfield" }
{ "_id" : "2", "Dept" : "SustainableLiving", "city":"Fairfield" }
```

10

- OUTLINE
- 🔍
- 

2. CRUD in MongoDB



3. CRUD in MongoDB



4. Searching an Array



5. Search with \$elemMatch



6. \$elemMatch



7. Searching an Object



8. Update Methods



9. Field Update Operators



10. Field Update Operators



# Field Update Operators

Update Operator	Description	Notes
<b>\$set</b>	Replaces the value of a field with the specified value	If the field does not exist, \$set will add a new field with the specified value
<b>\$inc</b>	Increments a field by a specified value, it accepts positive and negative values	If the field does not exist, \$inc creates the field and sets the field to the specified value
<b>\$unset</b>	Deletes a particular field, The specified value in the \$unset expression does not impact the operation.	If the field does not exist, then \$unset does nothing

11

- OUTLINE
- Search...


3. CRUD in MongoDB


4. Searching an Array


5. Search with \$elemMatch


6. \$elemMatch


7. Searching an Object


8. Update Methods


9. Field Update Operators


10. Field Update Operators


11. Field Update Operators



# Examples - Field Update Operators

```
{
  "_id" : "1", "students" : 250
}
db.col.update({_id:"1"}, { $inc : { "students":1, "exams":1 } })
{
  "_id" : "1", "students" : 251, "exams":1
}
```

```
{
  "_id" : "1", "students" : 250, "dept": "ComPro"
}
db.col.update({_id:"1"}, { $unset : { "dept":1 } })
{
  "_id" : "1", "students" : 250
}
```

## OUTLINE

Search...



4. Searching an Array



5. Search with  
\$elemMatch



6. \$elemMatch



7. Searching an Object



8. Update Methods



9. Field Update  
Operators



10. Field Update  
Operators



11. Field Update  
Operators



12. Examples - Field  
Update Operators





# Array Update Operators

```

//Original Document { "_id" : 1, "a" : [1, 2, 3, 4] }
db.testCol.update({_id:1}, { $set : { "a.2":5 } })
// output: { "_id" : 1, "a" : [1, 2, 5, 4] }
db.col.update({_id:1}, { $push : { "a": 6 } })
// output: { "_id" : 1, "a" : [1, 2, 5, 4, 6] }
db.col.update({_id:1}, { $pop : { "a": 1 } })
// output: { "_id" : 1, "a" : [1, 2, 5, 4] }
db.col.update({_id:1}, { $pop : { "a": -1 } })
// output: { "_id" : 1, "a" : [2, 5, 4] }
db.col.update({_id:1}, { $pushAll : { "a": [7, 8, 9] } })
// output: { "_id" : 1, "a" : [2, 5, 4, 7, 8, 9] }
db.col.update({_id:1}, { $pull : { "a": 5 } })
// output: { "_id" : 1, "a" : [2, 4, 7, 8, 9] }
db.col.update({_id:1}, { $pullAll : { "a": [2, 4, 8] } })
// output: { "_id" : 1, "a" : [7, 9] }
db.col.update({_id:1}, { $addToSet : { "a": 5 } })
// output: { "_id" : 1, "a" : [7, 9, 5] }
db.col.update({_id:1}, { $addToSet : { "a": 5 } })
// output: { "_id" : 1, "a" : [7, 9, 5] }
  
```

13

## OUTLINE



5. Search with \$elemMatch

6. \$elemMatch

7. Searching an Object

8. Update Methods

9. Field Update Operators

10. Field Update Operators

11. Field Update Operators

12. Examples - Field Update Operators

13. Array Update Operators





# \$ (The Array POSITION that matched)

The positional **\$** operator identifies an **element** in an array to **update** without explicitly specifying the position of the element in the array.

```
{
  "_id" : 1, "grades" : [ 85, 80, 84 ] }
db.students.updateOne( { _id: 1, grades: 80 }, { $set: { "grades.$" : 82 } } )
// { "_id" : 1, "grades" : [ 85, 82, 84 ] }
```

```
{
  _id: 4, grades: [ { total: 80, mean: 75, std: 8 },
                    { total: 85, mean: 90, std: 5 },
                    { total: 86, mean: 85, std: 8 } ] }
db.students.updateOne( { _id: 4, "grades.total": 85 },
                      { $set: { "grades.$.std" : 6 } } )
// { "total" : 85, "mean" : 90, "std" : 6 }
```

You must include the array field as part of the query.

14

## OUTLINE

Search...



6. \$elemMatch

7. Searching an Object

8. Update Methods

9. Field Update Operators

10. Field Update Operators

11. Field Update Operators

12. Examples - Field Update Operators

13. Array Update Operators

14. \$ (The Array POSITION that matched)





# \$elemMatch vs. arrayFilters

```

{ _id: 5, grades: [ { total: 80, mean: 75, std: 8 },
                    { total: 85, mean: 90, std: 5 },
                    { total: 90, mean: 85, std: 3 } ] }

db.students.updateOne(
  { _id: 5, grades: { $elemMatch: { total: { $lte: 90 }, mean: { $gt: 80 } } } },
  { $set: { "grades.$.std" : 6 } }
)

```

{ total: 85, mean: 90, std: 6 }  
 Only updates one element that matches our condition

```

db.students.updateOne(
  { _id: 5 },
  { $set: { "grades.$[obj].std" : 6 } },
  { arrayFilters: [{ "obj.total": { $lte: 90 }, "obj.mean": { $gt: 80 } }] }
)

```

{ total: 85, mean: 90, std: 6 }, { total: 90, mean: 85, std: 6 }  
 Always updates all elements in the array that match our condition

15

## OUTLINE



7. Searching an Object



8. Update Methods



9. Field Update Operators



10. Field Update Operators



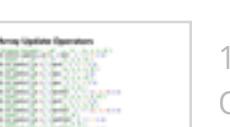
11. Field Update Operators



12. Examples - Field Update Operators



13. Array Update Operators



14. \$ (The Array POSITION that matched)



15. \$elemMatch vs. arrayFilters





# Using arrayFilters

```

{
  "_id" : 1,
  "grades" : [
    {
      "type": "quiz", "questions": [ 10, 8, 5 ] },
    {
      "type": "quiz", "questions": [ 8, 9, 6 ] },
    {
      "type": "hw", "questions": [ 5, 4, 3 ] },
    {
      "type": "exam", "questions": [ 25, 10, 23, 0 ] }
  ]
}

db.students.update(
  {},
  {
    $inc: { "grades.$[t].questions.$[score]": 2 },
    arrayFilters: [ { "t.type": "quiz" }, { "score": { $gte: 8 } } ],
    multi: true
  }
)

{
  "type": "quiz", "questions": [ 12, 10, 5 ] },
{
  "type": "quiz", "questions": [ 10, 11, 6 ] },

```

16

## OUTLINE



9. Field Update Operators

10. Field Update Operators

11. Field Update Operators

12. Examples - Field Update Operators

13. Array Update Operators

14. \$ (The Array POSITION that matched)

15. \$elemMatch vs. arrayFilters

16. Using arrayFilters

17. elemMatch vs arrayFilters





# elemMatch vs arrayFilters

```
{grades:[10,20,20]}
{grades:[10,20,20]}
```

```
db.test.update({grades: { $elemMatch: { $gt: 10 , $lt: 30 } } },
  {$set:{'grades.$': 15}},
  {multi: true})
```

```
{ "grades" : [ 10, 15, 20 ] }
{ "grades" : [ 10, 15, 20 ] }
```

```
db.test.update({},
  {$set:{'grades.$[c]': 30}},
  { arrayFilters: [ { "c": { $gt: 10 , $lt: 30 } } ], multi: true})
```

```
{ "grades" : [ 10, 30, 30 ] }
{ "grades" : [ 10, 30, 30 ] }
```

## OUTLINE

Search...



9. Field Update Operators



10. Field Update Operators



11. Field Update Operators



12. Examples - Field Update Operators



13. Array Update Operators



14. \$ (The Array POSITION that matched)



15. \$elemMatch vs. arrayFilters



16. Using arrayFilters



17. elemMatch vs arrayFilters





# Multi-updates

What really happens when we do multi-documents write operations is that inside MongoDB there is a **single thread** for every operation being executed (run sequentially inside single thread). Every write operation that affects more than one document is carefully coded in multi-tasking way to occasionally **yield** control to allow other operations to work on the same dataset.

- However, MongoDB guarantees an **individual doc** to be **atomic**, so no other R/W operation can get a half updated document.
- Multi-doc update operations are **not isolated transactions**. Check [\*\*`\$isolated\(\)`\*\*](#)

18

OUTLINE
<input type="text" value="Search..."/> <span>Search icon</span>
10. Field Update Operators
11. Field Update Operators
12. Examples - Field Update Operators
13. Array Update Operators
14. \$ (The Array POSITION that matched)
15. \$elemMatch vs. arrayFilters
16. Using arrayFilters
17. elemMatch vs arrayFilters
18. Multi-updates



# Performance

When we talk about performance we think adding memory, replacing the hard drive with SSD or using more powerful CPU (vertical scaling).

But what really affects performance is our algorithm and the logic we use in structuring our DB!

Understanding indexes leads to a better code and performance.



When would MongoDB use an Index?

25

## OUTLINE

Search...



11. Field Update Operators



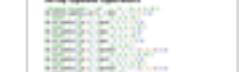
11. Field Update Operators

12. Examples - Field Update Operators



12. Examples - Field Update Operators

13. Array Update Operators



13. Array Update Operators

14. \$ (The Array POSITION that matched)



14. \$ (The Array POSITION that matched)

15. \$elemMatch vs. arrayFilters



15. \$elemMatch vs. arrayFilters

16. Using arrayFilters



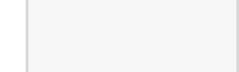
16. Using arrayFilters

17. elemMatch vs arrayFilters



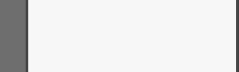
17. elemMatch vs arrayFilters

18. Multi-updates



18. Multi-updates

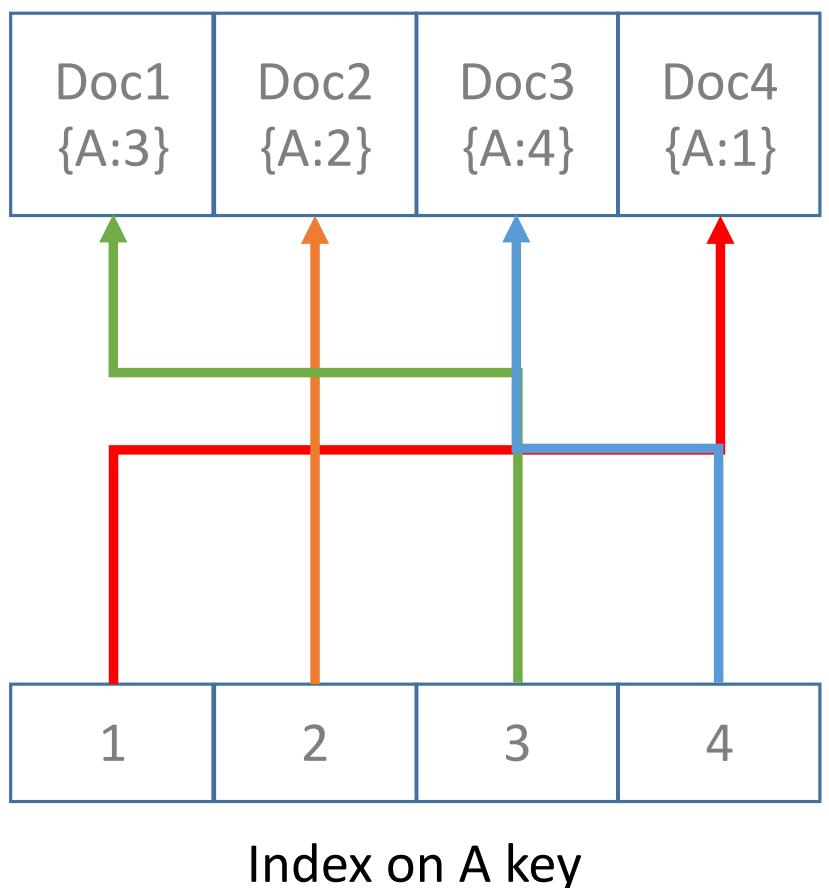
19. Performance





# Regular Index

- When we call `find()` a **full table/collection scan** will happen because our collection is not sorted which leads to slow performance.
- That's why we create an Index to boost the performance of our operations (*find, update, sort.. etc*)



26

## OUTLINE



12. Examples - Field Update Operators

13. Array Update Operators

14. \$ (The Array POSITION that matched)

15. \$elemMatch vs. arrayFilters

16. Using arrayFilters

17. elemMatch vs arrayFilters

18. Multi-updates

19. Performance

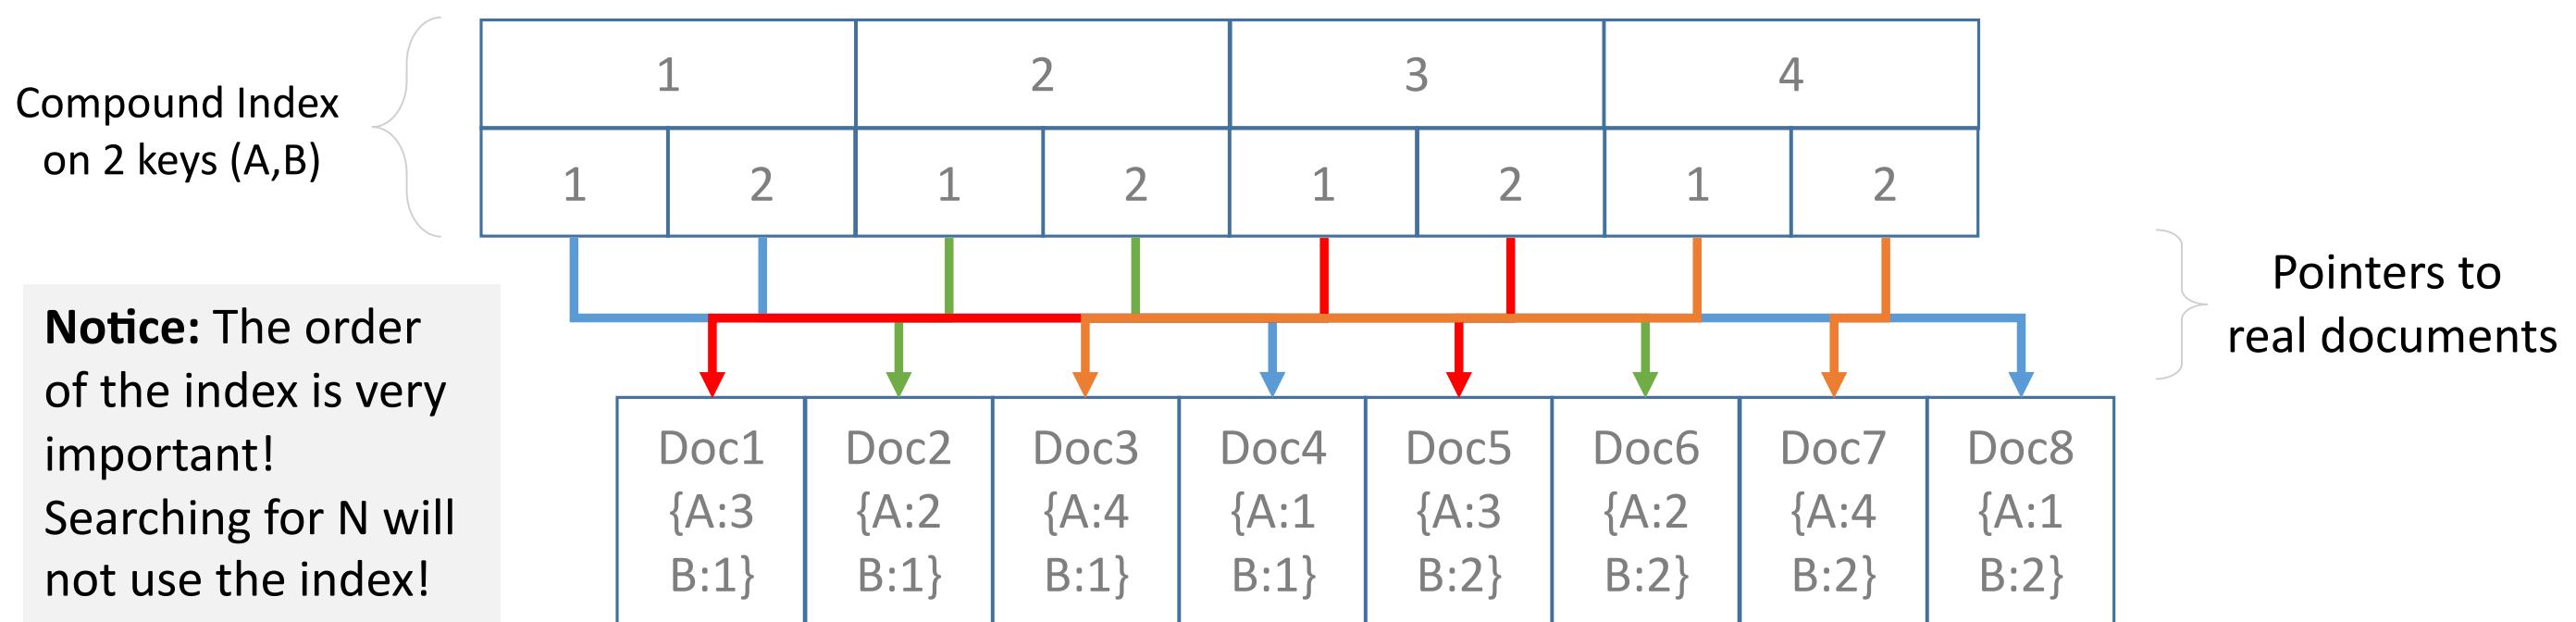
20. Regular Index





# Compound Index

When the index is composed of more than one key.



**Notice:** The index will be updated for every `insert()` operation (slow). Though reading from the index is very fast.

27

## OUTLINE

Search...



13. Array Update Operators



14. \$ (The Array POSITION that matched)



15. \$elemMatch vs. arrayFilters



16. Using arrayFilters



17. elemMatch vs arrayFilters



18. Multi-updates



19. Performance



20. Regular Index



21. Compound Index

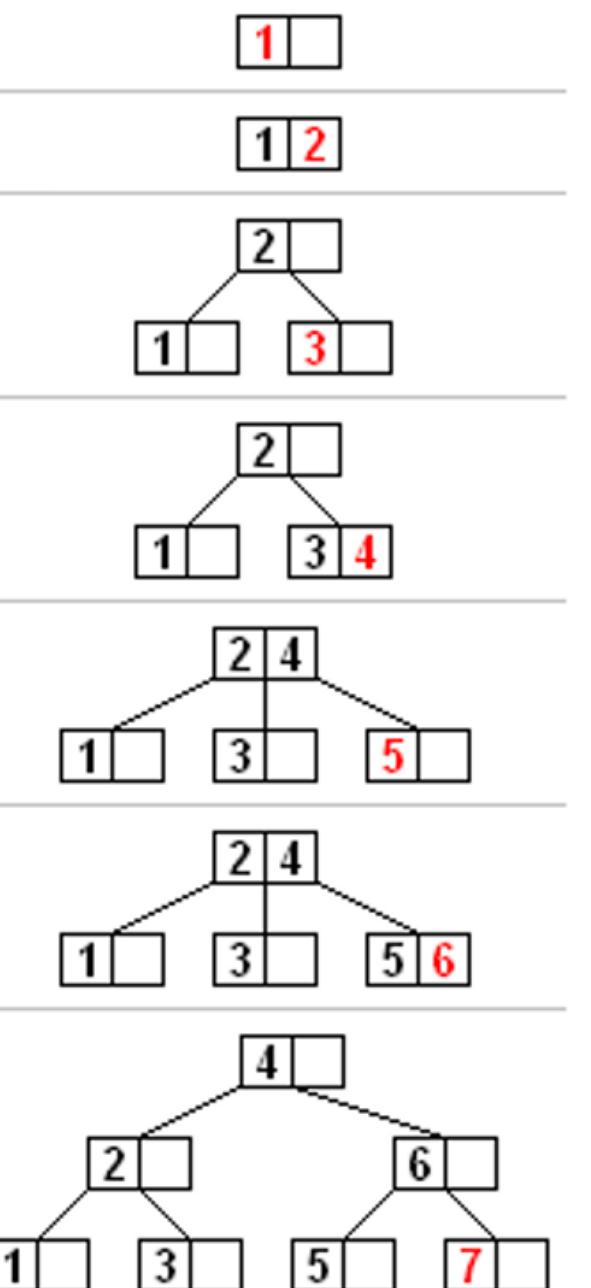




# B-Tree Algorithm

MongoDB uses B-Tree algorithm to maintain indexes. B-tree is a self-balancing tree data structure that keeps data sorted and allows searches, sequential access, insertions, and deletions in logarithmic time.

Algorithm	Average	Worst Case
Search	$O(\log n)$	$O(\log n)$
Insert	$O(\log n)$	$O(\log n)$
Delete	$O(\log n)$	$O(\log n)$



28

## OUTLINE



15. \$elemMatch vs. arrayFilters

16. Using arrayFilters

17. elemMatch vs arrayFilters

18. Multi-updates

19. Performance

20. Regular Index

21. Compound Index

22. B-Tree Algorithm

23. Index Operations





# Index Operations

```
// create a compound index on Key1 ASC, Key2 DESC
db.colName.createIndex({key1:1, key2:-1})

// show a list of all indexes in the collection
db.colName.getIndexes()

// drop the index
db.colName.dropIndex({key1:1, key2:-1})

// show useful information about index size
db.colName.stats()
```

**Note:** Creating an index will take time and space!

29

## OUTLINE



15. \$elemMatch vs. arrayFilters



16. Using arrayFilters



17. elemMatch vs arrayFilters



18. Multi-updates



19. Performance



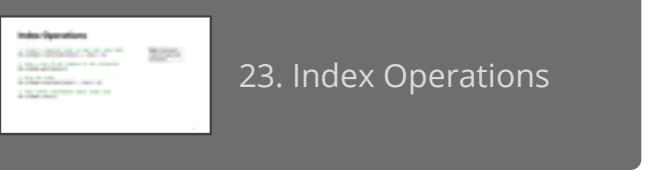
20. Regular Index



21. Compound Index



22. B-Tree Algorithm



23. Index Operations





# MultiKey Index

MongoDB supports index for **Array typed keys**. Only one Array typed key is allowed and MongoDB will create a compound key for all value combinations.

```
{
  name: '',
  sports: ['cycling', 'swimming'],
  address: [ {type: 'home', location: ['IA', 'Fairfield']},
             {type: 'work', location: ['NY', 'New York']} ]
}
db.colName.createIndex({name:1, sports:1})
db.colName.createIndex({name:1, address:1})
db.colName.createIndex({'address.type':1}) // supports rich documents
db.colName.createIndex({'address.location':1})
db.colName.createIndex({sports:1, address:1}) // will fail
```

30

## OUTLINE



16. Using arrayFilters



17. elemMatch vs arrayFilters



18. Multi-updates



19. Performance



20. Regular Index



21. Compound Index



22. B-Tree Algorithm



23. Index Operations



24. MultiKey Index





# Examples – MultiKey Index & Covered Index

```

db.col.insert({name: 'Asaad' , email: 'asaad@mum.edu' , work: 'MUM'})
db.col.createIndex({name:1, email:1}) // Create Compound Index
db.col.find({name: 'Asaad'}).explain() // Not Covered
db.col.find({name: 'Asaad'}, {name:1, email:1, _id:0}).explain() // Covered Index

db.col.insert({name: ['Mike', 'Mada'], email:'email@mum.edu'})
db.col.find({name: 'Mike'}).explain() // MultiKey Index

// Will fail: cannot index parallel arrays
db.col.insert({name: ['George', 'Angel'], email:['george@mum.edu', 'angel@mum.edu']})

```

**Covered Index:** The index is enough to return the results, there is no need to go to the original document

31

- OUTLINE
- 🔍
17. elemMatch vs arrayFilters

18. Multi-updates

19. Performance

20. Regular Index

21. Compound Index

22. B-Tree Algorithm

23. Index Operations

24. MultiKey Index

25. Examples – MultiKey Index & Covered Index
- 25 / 36
- 00:03 / 00:03
- 
- < PREV
- NEXT >



# Unique Index

```
db.col.insert({a:1, b:1})
db.col.insert({a:2, b:2})

// create a unique index on "a"
db.col.createIndex({a:1},{unique: true})

db.col.insert({a:2, b:2}) // will fail
```

## OUTLINE



18. Multi-updates



19. Performance



20. Regular Index



21. Compound Index



22. B-Tree Algorithm



23. Index Operations



24. MultiKey Index



25. Examples – MultiKey Index & Covered Index



26. Unique Index





# Sparse Index

```

db.col.insert({email:'asaad@mum.edu'})
db.col.insert({email:'mada@mum.edu'})
db.col.insert({name:'Mike', email:'mike@mum.edu'})
// create an index on "name"
db.col.createIndex({name:1}) // will work
// create a unique index on "name"
db.col.createIndex({name:1},{unique: true}) // will fail duplicate "null"
// create a unique index that doesn't have reference to the first two documents!
db.col.createIndex({name:1},{unique: true, sparse: true})

db.col.find({name:{$exists: false}}) // will use BasicCursor and return 2 docs
db.col.find({name:{$exists: false}}).hint({name:1}) // will use "name" index and
return nothing
  
```

**Sparse indexes** only contain entries for documents that have the indexed field. The index skips over any document that is missing the indexed field. The index is “sparse” because it does not include all documents of a collection.

33

## OUTLINE



 19. Performance

 20. Regular Index

 21. Compound Index

 22. B-Tree Algorithm

 23. Index Operations

 24. MultiKey Index

 25. Examples – MultiKey Index & Covered Index

 26. Unique Index

 27. Sparse Index





# Create Index in Production!

Creating an index for a large collection is a process that will block all other IO to the collection. MongoDB creates indexes in the foreground by default, despite it's a faster option, blocking other IO might not be an option for systems in production.

Creating an index in the background (writes can happen concurrently) is more efficient, but it's way more slower than foreground process.

```
db.collection.createIndex( { a: 1 }, { background: true } )
```

Is there another way to create an index fast in production?

34

- OUTLINE**
- 🔍
- 
20. Regular Index


21. Compound Index


22. B-Tree Algorithm


23. Index Operations


24. MultiKey Index


25. Examples – MultiKey Index & Covered Index


26. Unique Index


27. Sparse Index


28. Create Index in Production!
- 28 / 36
- 00:03 / 00:03
- 
- < PREV
- NEXT >



# How MongoDB chooses which index to use?

Let's say we have 3 indexes for our collection, the first time we perform a query, MongoDB will run three separate instances of the query using all three indexes, the first one finishes will be saved as the fastest index for this kind of queries. When data changes, MongoDB will run such a test again.

Remember we can always hint MongoDB to use a certain index:

```
// in Mongo Shell
db.test.find({name: 'Asaad'}).hint({name:1})
```

```
// in Node.js
db.collection('test').find({name: 'Asaad'}, {'hint': {'name':1}})
```

35

## OUTLINE



21. Compound Index



22. B-Tree Algorithm



23. Index Operations



24. MultiKey Index



25. Examples – MultiKey Index & Covered Index



26. Unique Index



27. Sparse Index



28. Create Index in Production!



29. How MongoDB chooses which index to use?





# Index Cardinality

For better performance, **it's important to keep indexes in memory**. Always check your DB stats `db.colName.stats()` and make sure you have enough memory for indexes.

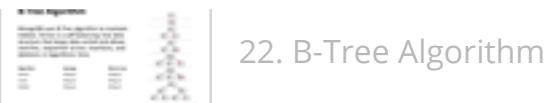
In proportion to collection size, index cardinality will be:

- **Regular Index** 1:1
- **Sparse Index** less or equal to documents size
- **MultiKey Index** significantly more than documents size

When you realize that you don't have enough resources for your index, it's time to consider Sharding.

36

## OUTLINE



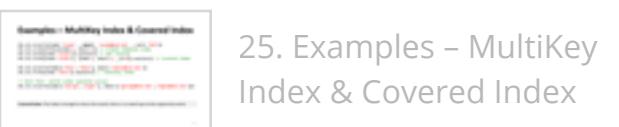
22. B-Tree Algorithm



23. Index Operations



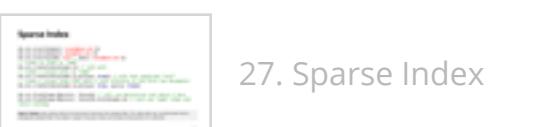
24. MultiKey Index



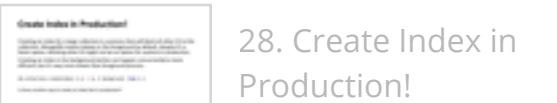
25. Examples – MultiKey Index &amp; Covered Index



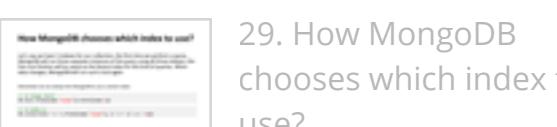
26. Unique Index



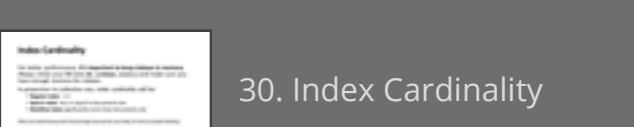
27. Sparse Index



28. Create Index in Production!



29. How MongoDB chooses which index to use?



30. Index Cardinality

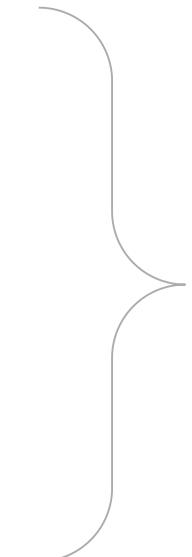




# Index Selectivity

It's important to understand our data before we select an index. Let's say we want to keep a log of activities in our application, **choosing an index with more varieties** will cut the time of searching and make our operations more efficient.

type	date
save	12 sep 2019
update	12 sep 2019
update	13 sep 2019
save	13 sep 2019
save	13 sep 2019



An index of (type, date) is not a good option as type is limited. While an index of (date, type) will improve the system performance because we can reach to the documents much faster

37

## OUTLINE



23. Index Operations



24. MultiKey Index



25. Examples – MultiKey Index &amp; Covered Index



26. Unique Index



27. Sparse Index



28. Create Index in Production!



29. How MongoDB chooses which index to use?



30. Index Cardinality



31. Index Selectivity





# FullText Index

When looking for a string value in MongoDB the full string will be checked against. While regular expressions could be a solution, it's very slow because it will search the entire collection. FullText Index is the best way to build an index for all words.

```
db.col.insert({mytext:'Hello goes to CS572 Class!'})
db.col.find({mytext:'Hello CS572 Class'}) // will find 0 docs
db.col.find({mytext:'CS572'}) // will find 0 docs
db.col.find({mytext: {$regex: 'CS572'}}) // will find 1 doc but slow!

db.col.createIndex({mytext: 'text'})
// will find our document fast using FullText index. spaces treated as OR
db.col.find({$text: {$search: 'Class CS572'}})
```

**Note:** MongoDB will log out all slow queries (more than 100ms) automatically to the shell. To write the log on the disk use a [Profiler](#)

38

## OUTLINE



24. MultiKey Index

25. Examples – MultiKey Index & Covered Index

26. Unique Index

27. Sparse Index

28. Create Index in Production!

29. How MongoDB chooses which index to use?

30. Index Cardinality

31. Index Selectivity

32. FullText Index





# A More Relevant FullText Search Results

FullText search will calculate a **textScore** value automatically for me, we can add this field to our results and sort them based on that field to get the documents that are more relevant.

```
{$text: {$search: 'text1 text2'} }
```

```
db.collection('colName').find({query}, { score: {$meta:'textScore'} } )  
    .sort( { score: {$meta:'textScore'} } )
```

The { \$meta: "textScore" } expression provides information on the processing of the \$text operation.

39

- OUTLINE**
- Search... 🔍
-  25. Examples – MultiKey Index & Covered Index
  -  26. Unique Index
  -  27. Sparse Index
  -  28. Create Index in Production!
  -  29. How MongoDB chooses which index to use?
  -  30. Index Cardinality
  -  31. Index Selectivity
  -  32. FullText Index
  -  33. A More Relevant FullText Search Results



# Distance Between 2 Points - The Haversine Formula

For two points on a sphere (of radius R) with latitudes  $\phi_1$  and  $\phi_2$ , latitude separation  $\Delta\phi = \phi_1 - \phi_2$ , and longitude separation  $\Delta\lambda$  the distance d between the two points

$$\text{haversin}\left(\frac{d}{R}\right) = \text{haversin}(\Delta\phi) + \cos(\phi_1) \cos(\phi_2) \text{haversin}(\Delta\lambda)$$

$$\text{haversin}(\theta) = \frac{\text{versin}(\theta)}{2} = \sin^2\left(\frac{\theta}{2}\right)$$

$$\text{versin}(\theta) = 1 - \cos(\theta) = 2 \sin^2\left(\frac{\theta}{2}\right)$$

## The Haversine Formula in MySQL

```
SELECT *, 3956 * 2 * ASIN(SQRT(POWER(SIN(@orig_lat - abs(dest.lat)) * pi()/180 / 2), 2) + COS(@orig_lat * pi()/180) * COS(abs(dest.lat) * pi()/180) * POWER(SIN(@orig_lon - dest.lon) * pi()/180 / 2), 2)) as distance
FROM hotels dest HAVING distance < @dist ORDER BY distance LIMIT 10;
```

40

## OUTLINE



26. Unique Index



27. Sparse Index



28. Create Index in Production!



29. How MongoDB chooses which index to use?



30. Index Cardinality



31. Index Selectivity



32. FullText Index



33. A More Relevant FullText Search Results



34. Distance Between 2 Points - The Haversine Formula





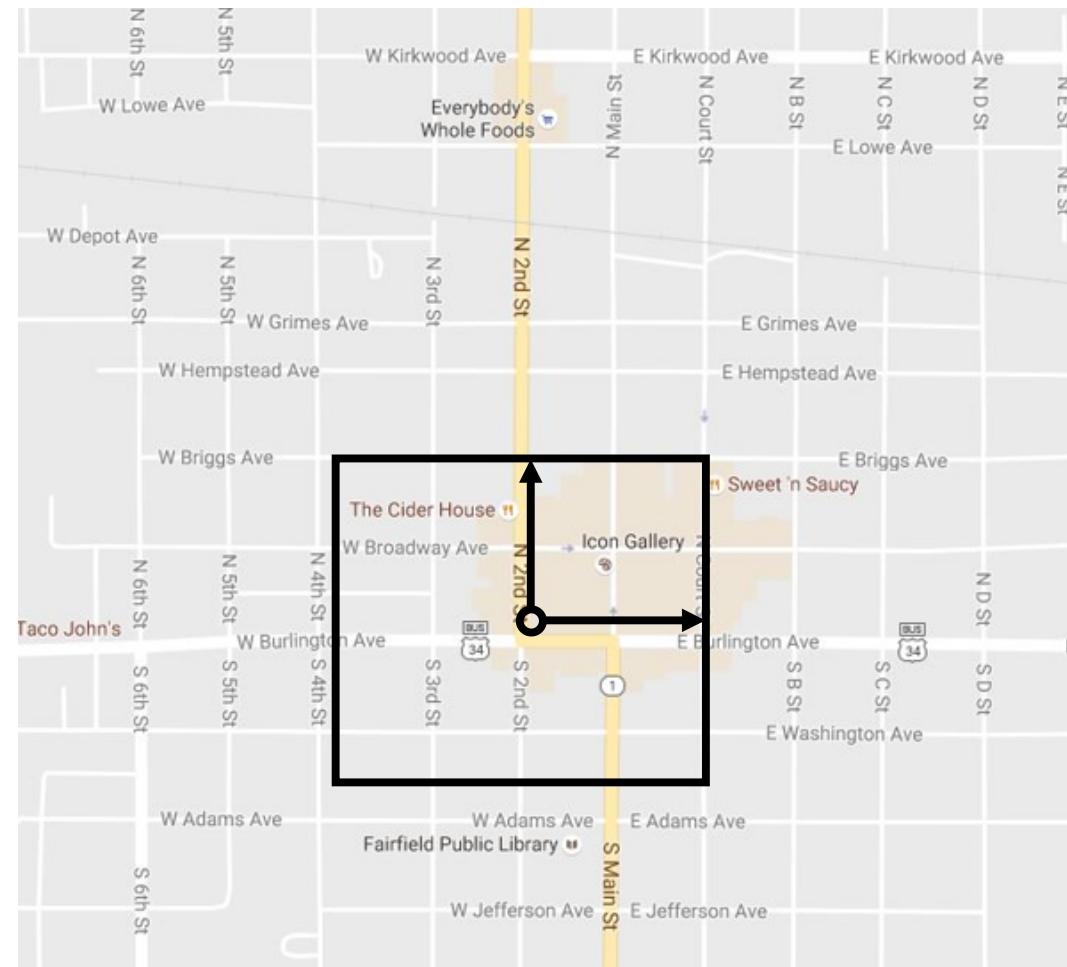
# Geospatial Indexes (2D)

To find nearest locations in MongoDB all we need is :

- Have all locations saved in each collection with [long, lat] array
- Create an Geospatial Index (2d) on that key.

**Code Example:**

```
db.collection.insert{name:'Restaurant', location: [Long, Lat]}
db.collection.createIndex({location: '2d'})
// results will be sorted by increasing distance (closest to furthest)
db.collection.find({location: {$near:[currentLong, currentLat]}}).limit(10)
```



41

## OUTLINE

Search...



27. Sparse Index

28. Create Index in Production!

29. How MongoDB chooses which index to use?

30. Index Cardinality

31. Index Selectivity

32. FullText Index

33. A More Relevant FullText Search Results

34. Distance Between 2 Points - The Haversine Formula

35. Geospatial Indexes (2D)





# Geospatial Spherical (3D)

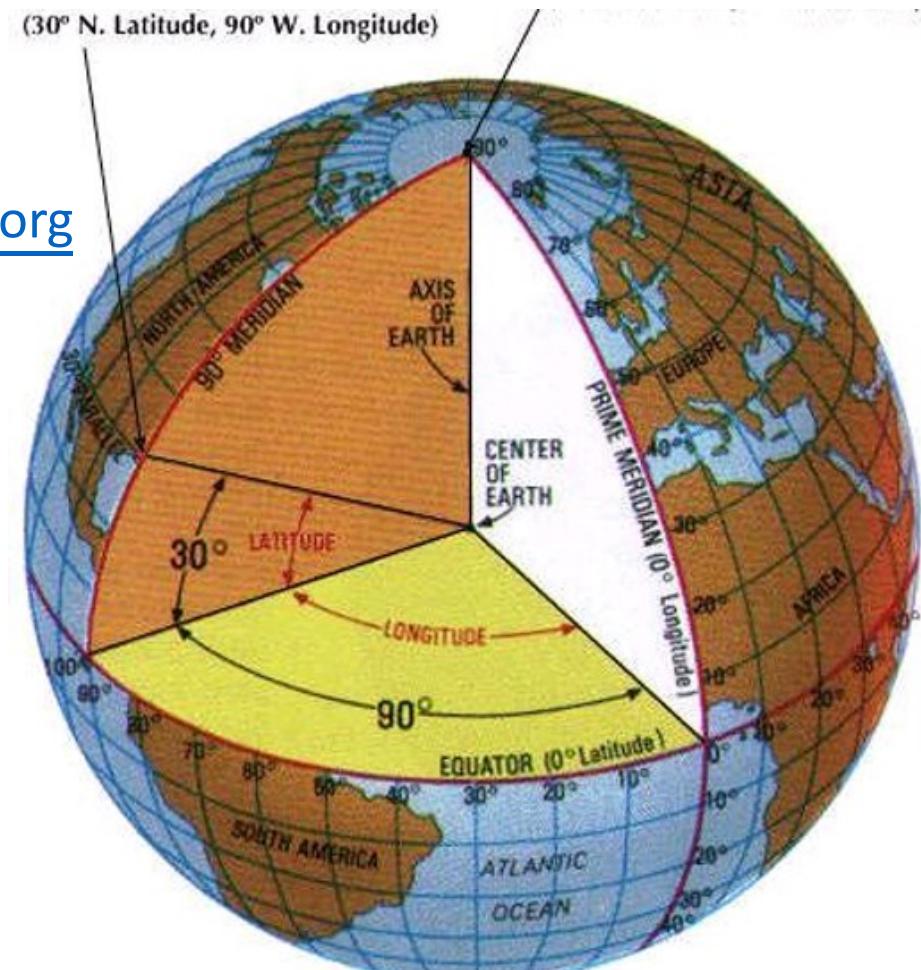
Searching on spherical surface is more accurate and gives better results.

You can check encoding of geographic data structures at <http://geojson.org>

```
{
  "_id": 1,
  "name": "MUM University",
  "city": "Fairfield",
  "location": {
    "type": "Point",
    "coordinates": [-91.9612747, 41.0132949]
  }
}
```

```
db.places.createIndex({'location': '2dsphere'})
db.places.find({location: {
  $near: {$geometry: {type: "Point", coordinates: [-91.9627755, 41.0085809]}, $maxDistance: 2000 } } })
```

In meters



42

## OUTLINE

Search...



28. Create Index in Production!



29. How MongoDB chooses which index to use?



30. Index Cardinality



31. Index Selectivity



32. FullText Index



33. A More Relevant FullText Search Results



34. Distance Between 2 Points - The Haversine Formula



35. Geospatial Indexes (2D)



36. Geospatial Spherical (3D)

